

Exercise 2

For this exercise , you will be working with the [Titanic Data Set from Kaggle \(https://www.kaggle.com/c/titanic\)](https://www.kaggle.com/c/titanic). This is a very famous data set and very often is a student's first step in Data Analytics!

The Dataset has been given to you on D2L. You need to download the .csv file from your assignment folder. The above link is just for a reference story about the data.

1- For this assignment, you need to perform exploratory data analysis and answer at least three hypotheses based on the dataset. You may need to use your knowledge of statistics to analyze this data.

Here are three possible hypotheses that you can define for this dataset (you can define your own hypotheses as well):

- Determine if the survival rate is associated to the class of passenger
- Determine if the survival rate is associated to the gender
- Determine the survival rate is associated to the age

2- For each hypothesis, you need to make at least one plot.

3- Write a summary of your findings in one page (e.g., summary statistics, plots) and submit the pdf file. Therefore, for part 2 of your assignment, you need to submit one jupyter notebook file and one pdf file.

This will be your first end to end data analysis project. For this assignment, you will be graded on your overall analysis, and your final report.

4- Push your code and project to github and provide the link to your code here.

Ensure that your github project is organized to at least couple of main folders, ensure that you have the README file as well:

- Src
- Data
- Docs
- Results

Read this link for further info: <https://gist.github.com/ericmjl/27e50331f24db3e8f957d1fe7bbbe510>
(<https://gist.github.com/ericmjl/27e50331f24db3e8f957d1fe7bbbe510>).

Topic will be covered

The framework of this notebook is:

1. Exploratory data analysis\
 - 1) NaN Values\
 - 2) Survived\
 - 3) Pclass\
 - 4) Name Column\
 - 5) Age Column\
 - 6) Sex Column
2. Hypothesis\

- 1) Determine if the survival rate is associated to the class of passenger\
- 2) Determine if the survival rate is associated to the gender\
- 3) Determine the survival rate is associated to the age

In [58]:

```
import pandas as pd
import numpy as np
import pickle

# visualization
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

# Enable Jupyter Notebook's intellisense
%config IPCompleter.greedy=True

# We want to see whole content (non-truncated)
pd.set_option('display.max_colwidth', None)
```

In [59]:

```
df = pd.read_csv("titanic.csv")
display(df.head())

print(df.info())
print(df.describe())
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500		S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)	female	38.0	1	0	PC 17599	71.2833		C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250		S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000		C
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500		S

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
Column Non-Null Count Dtype

0 PassengerId 891 non-null int64
1 Survived 891 non-null int64
2 Pclass 891 non-null int64
3 Name 891 non-null object
4 Sex 891 non-null object
5 Age 714 non-null float64
6 SibSp 891 non-null int64
7 Parch 891 non-null int64
8 Ticket 891 non-null object
9 Fare 891 non-null float64
10 Cabin 204 non-null object
11 Embarked 889 non-null object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
None

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	714.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	14.526497	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	
25%	223.500000	0.000000	2.000000	20.125000	0.000000	
50%	446.000000	0.000000	3.000000	28.000000	0.000000	

75%	668.500000	1.000000	3.000000	38.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

Exploratory Data Analysis

NaN values

In [60]:

```
print(df.isnull().sum())
```

```

PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age             177
SibSp            0
Parch            0
Ticket           0
Fare             0
Cabin           687
Embarked         2
dtype: int64

```

From the records, we can see that 4 features have missing values:

Age: Age is fractional if less than 1. -> Numerical Variable

Cabin: Cabin number -> Categorical variable

Embarked: Port of Embarkation (3 categories) -> Categorical variable

Fare: Numerical: Only 1 missing in the test dataset. Can be replace by the mean in the training set.

Survived

In [61]:

```
#COUNT THE SURVIVOIRS AND UNSURVIVOIRS
survivours = df.Survived.sum()
unsurvivours = len(df) - survivours


print(survivours, "people survived")
print(unsurvivours, "people didn't survive")

#Visualize Survived column. (1=Survived, 0=Not Survived)
sns.countplot(x="Survived", data=df)
plt.show()

#PRINT PROPORTIONS
print(df["Survived"].value_counts(normalize=True))
```

342 people survived

549 people didn't survive



0 0.616162

1 0.383838

Name: Survived, dtype: float64

Pclass

Pclass column contains the status of the passengers.

1 = Upper

2 = Middle


3 = Lower

In [62]:

```
#hue parameter represents which column in the data frame, you want to use for color encoding
sns.countplot(x="Survived", hue="Pclass", data=df)
```

Out[62]:

<AxesSubplot:xlabel='Survived', ylabel='count'>



Name Column

In [63]:

```
# Display first five rows of the Name column
display(df[["Name"]].head())
```

	Name
0	Braund, Mr. Owen Harris
1	Cumings, Mrs. John Bradley (Florence Briggs Thayer)
2	Heikkinen, Miss. Laina
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)
4	Allen, Mr. William Henry

Age Column

In [64]:


```
# Print the missing values in Age column
print(df["Age"].isnull().sum())
```

In [65]:

```
# There are 177 missing values in Age column. Now, Let's Look at the distribution of ages b

# Survived by age
sns.histplot(df[df.Survived==1]["Age"],color="y", bins=7, label="1")

# Death by age
sns.histplot(df[df.Survived==0]["Age"], bins=7, label="0")
plt.legend()
plt.title("Age Distribution")
plt.show()
```




Sex Column

In [66]:

```
# Visualize with a countplot
sns.countplot(x="Sex", hue="Survived", data=df)
plt.show()

# Proportion of people survived for each class
print(df["Survived"].groupby(df["Sex"]).mean())

# How many people we have in each class?
print(df["Sex"].value_counts())
```



```
Sex
female    0.742038
male      0.188908
Name: Survived, dtype: float64
male      577
female    314
Name: Sex, dtype: int64
```

Hypothesis

1. Determine if the survival rate is associated to the class of passenger.

Let's check how the passenger class (Pclass) impacted the survival of the passengers.

In [67]:

```
total_passenger_count = len(df['Pclass'])
first_count = (df['Pclass'] == 1).sum()
second_count = (df['Pclass'] == 2).sum()
third_count = (df['Pclass'] == 3).sum()

per_first = first_count / total_passenger_count * 100
per_second = second_count / total_passenger_count * 100
per_third = third_count / total_passenger_count * 100
```


In [68]:

```
print('First class percentage = ', per_first_class, ' Count: ', first_class_count)
print('Second class percentage = ', per_second_class, ' Count: ', second_class_count)
print('Third class percentage = ', per_third_class, ' Count: ', third_class_count)
```

```
First class percentage = 24.242424242424242 Count: 216
Second class percentage = 20.65095398428732 Count: 184
Third class percentage = 55.106621773288445 Count: 491
```

With these survival rates, we can see that:

The first class represents $\approx 24\%$ of the passengers, but $\approx 40\%$ of the survivors

The second class represents $\approx 20\%$ of the passengers, but $\approx 25\%$ of the survivors

The third class represents $\approx 55\%$ of the passengers, but $\approx 34\%$ of the survivors

In [69]:

```
df[['Pclass', 'Survived']].groupby(['Pclass'], as_index=False).mean()
```

Out[69]:


	Pclass	Survived
0	1	0.629630
1	2	0.472826
2	3	0.242363

In [72]:

```
# Visualize with a countplot (1=Survived, 0=Not Survived)
sns.countplot(x="Pclass", hue="Survived", data=df)
plt.show()

# Proportion of people survived for each class
print(df["Survived"].groupby(df["Pclass"]).mean())

# How many people we have in each class?
print(df["Pclass"].value_counts())
```



```
Pclass
1    0.629630
2    0.472826
3    0.242363
Name: Survived, dtype: float64


3    491
1    216
2    184
Name: Pclass, dtype: int64
```

In [73]:

```
#df.groupby('Pclass').Survived.mean().plot(kind='bar')
sns.barplot(x='Pclass', y='Survived', data=df)
```

Out[73]:

<AxesSubplot:xlabel='Pclass', ylabel='Survived'>



From the above two graph we can see that higher class passengers have better survival chance. Therefore, we can say that the hypothesis is correct, survival rate is associated to the class of passenger.

2.Determine if the survival rate is associated to the gender

Now we will analyze if there is any relation between gender and survival rate. First, we'll get the count of survivals and deaths by gender.

In [74]:

```
df[['Sex', 'Survived']].groupby(['Sex'], as_index=False).mean()
```

Out[74]:


	Sex	Survived
0	female	0.742038
1	male	0.188908

In [41]:

```
# Visualize with a countplot
sns.countplot(x="Sex", hue="Survived", data=df)
plt.show()

# Proportion of people survived for each class
print(df["Survived"].groupby(df["Sex"]).mean())

# How many people we have in each class?
print(df["Sex"].value_counts())
```



```
Sex
female    0.742038
male      0.188908
Name: Survived, dtype: float64


Sex
male      577
female    314
Name: Sex, dtype: int64
```

In [44]:

```
#df.groupby('Sex').Survived.mean().plot(kind='bar')  
sns.barplot(x='Sex', y='Survived', data=df)
```

Out[44]:

<AxesSubplot:xlabel='Sex', ylabel='Survived'>



From the above two graph we can see that Females had better survival chance. Therefore, we can say that the survival rate is associated to the gender of the passenger.


3. Determine the survival rate is associated to the age

In [12]:

```
#Determine the survival rate is associated to the age  
sns.displot(df['Age'].dropna(), kde=False, bins=30)
```

Out[12]:

<seaborn.axisgrid.FacetGrid at 0x2b44fbbe910>



In [46]:

```
df_age = df[['Age' , 'Survived']].dropna(how='any')  
df_age['Age'] = (np.floor(df_age['Age'])).astype(int)  
df_age.shape
```

Out[46]:

(714, 2)

In [47]:

```
print('Youngest passenger: ' + str(df_age['Age'].min() * 12) + ' months old')
print('\nOldest passenger: ' + str(df_age['Age'].max()) + ' years old')
```

Youngest passenger: 0 months old

Oldest passenger: 80 years old

We rounded the ages up for the sake of this analysis, so the youngest passenger on board was a few months old, and the oldest was 80 years old.

Now let's order our data and create a new DataFrame `df_ages_survival`. It will be Age indexed. For each age, 4 columns will give us the following information:

Number of passengers of index age that survived

Number of passengers of index age that died

Total number of passengers of index age

Percentage of survivors of that index age

In [48]:

```
ages_list = df_age['Age'].unique()
ages_list.sort()
ages_list
```

Out[48]:

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
        34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50,
        51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 70,
        71, 74, 80])
```

In [49]:

```
df_ages_survival = pd.DataFrame(index=ages_list, columns=['Survived', 'Deaths', 'Total', 'Percentage'])

df_ages_survival['Survived'] = df_age.groupby('Age')['Survived'].sum()
df_ages_survival['Total'] = df_age.groupby('Age').count()
df_ages_survival['Deaths'] = df_ages_survival['Total'] - df_ages_survival['Survived']
df_ages_survival['Percentage'] = round(df_age.groupby('Age')['Survived'].mean() * 100, 2)

df_ages_survival.head()
```


Out[49]:

	Survived	Deaths	Total	Percentage
0	7	0	7	100.00
1	5	2	7	71.43
2	3	7	10	30.00
3	5	1	6	83.33
4	7	3	10	70.00

Let's take a look at a scatter plot of our data, and see if we can identify any outliers.

In [50]:

```
x = df_ages_survival['Percentage'].index
y = df_ages_survival['Percentage']
plt.scatter(x, y)
m, b = np.polyfit(x, y, 1)
plt.plot(x, y, '.')
plt.title('Survival rates by age')
plt.xlabel('Age in years')
plt.ylabel('Percentage of survivors')
plt.show()
```




We can identify a tendency where the lower y values have high x coordinates, and high y values have low x coordinates. This means that apparently, the younger you are, the more chances you have to survive.

However, there are three dots on the upper right side of the graph that seem to be outliers. They are not representative of the general tendency: let's identify them, remove them from the DataFrame, plot it again and fit a regression line.

In [51]:

```
df_ages_survival = df_ages_survival.drop(df_ages_survival[(df_ages_survival['Percentage'] =
x = df_ages_survival['Percentage'].index
y = df_ages_survival['Percentage']
plt.scatter(x, y)
m, b = np.polyfit(x, y, 1)
plt.plot(x, y, '.')
plt.plot(x, m*x + b, '-')
plt.title('Survival rates by age')
plt.xlabel('Age in years')
plt.ylabel('Percentage of survivors')
plt.show()
```




In []:

The regression line clearly shows that the younger you were, the higher your chances to sur

In [52]:

```
# Survived by age
sns.histplot(df[df.Survived==1]["Age"],color="y", bins=7, label="1")

# Death by age
sns.histplot(df[df.Survived==0]["Age"], bins=7, label="0")
plt.legend()
plt.title("Age Distribution")
plt.show()
```



We can definitely see a spike in the distribution of survival passengers when the age is small, indicating children had a higher survival rate.

It would seem that children up to 10 years old had a higher chance of survival than people over 10 years old.

From the above graphs we can say that the younger you were, the higher your chances to survive. Therefore, we can say that the survival rate is associated with the age of the passenger.